Analysis of Lead Gender and Box Office

by Max Woolf (@minimaxir)

This notebook is licensed under the MIT License. If you use the code or data visualization designs contained within this notebook, it would be greatly appreciated if proper attribution is given back to this notebook and/or myself. Thanks! :)

```
options(warn = -1)

# IMPORTANT: This assumes that all packages in "Rstart.R" are installed,
# and the fonts "Source Sans Pro" and "Open Sans Condensed Bold" are installed
# via extrafont. If ggplot2 charts fail to render, you may need to change/remove the theme
    call.

source("Rstart.R")
library(outliers)

sessionInfo()
```

```
Attaching package: ‘’dplyr

The following objects are masked from ‘package:’stats:

    filter, lag

The following objects are masked from ‘package:’base:

    intersect, setdiff, setequal, union

Registering fonts with R

Attaching package: ‘’scales

The following objects are masked from ‘package:’readr:

    col_factor, col_numeric









R version 3.2.3 (2015-12-10)
Platform: x86_64-apple-darwin13.4.0 (64-bit)
Running under: OS X 10.11.4 (El Capitan)

locale:
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8

attached base packages:
[1] grid      stats     graphics  grDevices utils     datasets  methods
[8] base

other attached packages:
```

```
[1] outliers_0.14      stringr_1.0.0      digest_0.6.8       RColorBrewer_1.1-2
[5] scales_0.3.0       extrafont_0.17     ggplot2_2.0.0      dplyr_0.4.3
[9] readr_0.1.1

loaded via a namespace (and not attached):
 [1] Rcpp_0.12.1      Rttf2pt1_1.3.3   magrittr_1.5     munsell_0.4.2
 [5] uuid_0.1-2       colorspace_1.2-6 R6_2.1.1         plyr_1.8.3
 [9] tools_3.2.3      parallel_3.2.3   gtable_0.1.2     DBI_0.3.1
[13] extrafontdb_1.0  assertthat_0.1   IRdisplay_0.3    repr_0.4
[17] base64enc_0.1-3  IRkernel_0.5     evaluate_0.8     rzmq_0.7.7
[21] stringi_0.5-5    jsonlite_0.9.19
```

Process the Data

Take the movies data, load in R friendly format, and combine with Rotten Tomatoes data.

```
df <- read_delim("~/Downloads/omdb0316/omdbMovies.txt", "\t",
    col_types="iccicccccccidi_c_____")
df_tomatoes <- read_delim("~/Downloads/omdb0316/tomatoes.txt",  "\t",
    col_types="i_diiiicidi_cc_c")
df <- df %>% left_join(df_tomatoes, by="ID")
rm(df_tomatoes)
```

```
|======================================================================================| 100%  435
    MB
```

```
parseBoxOffice <- function(x) {
    unit <- 0
    if (is.na(x) | x=="") {return (NA)}
    if (substr(x, nchar(x), nchar(x)) == "k") {unit <- 10^3}
    else {unit <- 10^6}

    number <- as.numeric(substr(x,2,nchar(x)-1))

    return(number * unit)
}

df <- df %>% mutate(BoxOffice = as.numeric(sapply(BoxOffice, parseBoxOffice)))
```

```
df_dup <- df %>% select(Title, Year) %>% mutate(Title = gsub("The ", "", Title))
dup <- duplicated(df_dup)   # find entry indices which are duplicates
rm(df_dup)   # remove temp dataframe

df <- df %>% filter(!dup)   # keep entries which are *not* dups
```

Inflation

```
inflation <- read_csv("http://research.stlouisfed.org/fred2/data/CPIAUCSL.csv") %>%
                    group_by(Year = as.integer(substr(DATE, 1, 4))) %>%
                    summarize(Avg_Value = mean(VALUE)) %>%   # average across all months
                    mutate(Adjust = tail(Avg_Value, 1) / Avg_Value)   # normalize by
                        most-recent year
```

```
df <- df %>% inner_join(inflation) %>% mutate(AdjBoxOffice = floor(BoxOffice * Adjust))
```

```
Joining by: "Year"
```

Select only data we need now.

```
df <- df %>% filter(Year >= 2000, AdjBoxOffice >= 10^7, Cast != '') %>%
        select(imdbID, Title, Year, Cast, Meter, Metacritic, AdjBoxOffice) %>%
        arrange(desc(AdjBoxOffice))

#write.csv(df, "test.csv", row.names=F)
print(nrow(df))
```

```
[1] 2048
```

Determine Gender of Lead

```
# Helper function to get first actor given a string of actors
getLeadActor <- function(actors) {
    return(unlist(strsplit(actors,", "))[1])
}

# Unicode issues during testing, so use string w/ unicode as a test case
print(getLeadActor("Will Smith, Robert De Niro, Renée Zellweger, Jack Black"))
```

```
[1] "Will Smith"
```

```
df$LeadActor <- as.character(lapply(enc2utf8(df$Cast), getLeadActor))

print(head(df %>% select(Title, LeadActor)))
print(head(df %>% filter(imdbID=="tt0307453") %>% select(Title, LeadActor)))
```

```
Source: local data frame [6 x 2]

                                     Title         LeadActor
                                     (chr)             (chr)
1 Star Wars: Episode VII - The Force Awakens     Harrison Ford
2                                   Avatar   Sam Worthington
3                            Jurassic World       Chris Pratt
4                              The Avengers Robert Downey Jr.
5                           The Dark Knight    Christian Bale
6                                  Shrek 2        Mike Myers
Source: local data frame [1 x 2]

       Title  LeadActor
       (chr)      (chr)
1 Shark Tale Will Smith
```

Attempt #1: Merge known gender data of actors. (via list from Matt Daniels)

```
actor_gender <- read_csv("actor_list.csv") %>% select(LeadActor = name, Gender = gender)

print(head(actor_gender))
```

```
Source: local data frame [6 x 2]

          LeadActor Gender
              (chr)  (chr)
1       Keir Dullea      m
2     Gary Lockwood      m
3 William Sylvester      m
4    Daniel Richter      m
5  Leonard Rossiter      m
6   Margaret Tyzack      f
```

```
df <- df %>% left_join(actor_gender)

print(head(df %>% select(Title, LeadActor, Gender)))
```

```
Joining by: "LeadActor"


Source: local data frame [6 x 3]

                                      Title         LeadActor Gender
                                      (chr)             (chr)  (chr)
1 Star Wars: Episode VII - The Force Awakens     Harrison Ford      m
2                                     Avatar   Sam Worthington     NA
3                              Jurassic World       Chris Pratt     NA
4                                The Avengers Robert Downey Jr.     NA
5                             The Dark Knight    Christian Bale     NA
6                                    Shrek 2        Mike Myers     NA
```

Attempt #2: Determine gender from most-likely guess from first name. (Using male and female lists from Carnegie-Mellon University)

```
male_names <- unlist(read_delim("male_names.txt", "\n", skip = 6, col_names=F))
female_names <- unlist(read_delim("female_names.txt", "\n", skip = 6, col_names=F))

print(head(male_names))
print(head(female_names))
```

```
     X11      X12      X13      X14      X15      X16
 "Aamir"  "Aaron"  "Abbey"  "Abbie"  "Abbot" "Abbott"
      X11       X12       X13       X14       X15       X16
"Abagael" "Abagail"    "Abbe"   "Abbey"    "Abbi"   "Abbie"
```

```
getGenderFromFullName <- function (full_name) {
    first_name <- unlist(strsplit(full_name, " "))[1]
    gender <- ifelse(first_name %in% male_names, "m",
                     ifelse(first_name %in% female_names, "f", "[EDIT ME]"))
    return (gender)
}

print(getGenderFromFullName("Sam Worthington"))
print(getGenderFromFullName("Kristen Wiig"))
```

```
[1] "m"
[1] "f"
```

```
gender_guess <- as.character(lapply(as.character(df$LeadActor), getGenderFromFullName))

# if a known gender from IMDB is present, use that; else, use the gender guess
df$Gender <- ifelse(is.na(df$Gender), gender_guess, df$Gender)

print(head(df %>% select(Title, LeadActor, Gender)))
print(tail(df %>% select(Title, LeadActor, Gender)))
```

```
Source: local data frame [6 x 3]

                                     Title         LeadActor Gender
                                     (chr)             (chr)  (chr)
1 Star Wars: Episode VII - The Force Awakens     Harrison Ford      m
2                                    Avatar   Sam Worthington      m
3                             Jurassic World       Chris Pratt      m
4                               The Avengers Robert Downey Jr.      m
5                            The Dark Knight    Christian Bale      m
6                                   Shrek 2        Mike Myers      m
Source: local data frame [6 x 3]

                        Title              LeadActor    Gender
                        (chr)                  (chr)     (chr)
1                     The Man      Samuel L. Jackson         m
2           Aliens of the Deep Anatoly M. Sagalevitch         m
3                 A Single Man            Colin Firth         m
4                      Pollock              Ed Harris         m
5              Connie and Carla           Nia Vardalos [EDIT ME]
6 I Don't Know How She Does It   Sarah Jessica Parker         f
```

Attempt #3: Manually edit edge cases in a GUI (not shown)

```
write.csv(df, "movie_gender_intermediate.csv", row.names=F)
```

Begin the Analysis

Reload the updated dataset (a few rows were removed due to being dupes)

```
df <- read_csv("movie_gender_fixed.csv")

print(head(df %>% select(Title, LeadActor, Gender)))
print(nrow(df))
```

```
Source: local data frame [6 x 3]

                                     Title         LeadActor Gender
                                     (chr)             (chr)  (chr)
1 Star Wars: Episode VII - The Force Awakens      Daisy Ridley      f
2                                    Avatar   Sam Worthington      m
3                             Jurassic World       Chris Pratt      m
4                               The Avengers Robert Downey Jr.      m
5                            The Dark Knight    Christian Bale      m
6                                   Shrek 2        Mike Myers      m
[1] 2020
```

Can we remove any points (e.g. Star Wars) as outliers? (tests via R Explorations)

```
AdjRevenue <- unlist(head(df %>% select(AdjBoxOffice)))

dixon.test(AdjRevenue, opposite=F)
grubbs.test(AdjRevenue, opposite=F)
chisq.out.test(AdjRevenue, variance=var(AdjRevenue), opposite=F)
```

```
    Dixon test for outliers

data:  AdjRevenue
Q.AdjBoxOffice1 = 0.23695, p-value = 0.8916
alternative hypothesis: highest value 934381231 is an outlier
```

```
    Grubbs test for one outlier

data:  AdjRevenue
G.AdjBoxOffice1 = 1.52510, U = 0.44175, p-value = 0.2634
alternative hypothesis: highest value 934381231 is an outlier
```

```
    chi-squared test for outlier

data:  AdjRevenue
X-squared.AdjBoxOffice1 = 2.326, p-value = 0.1272
alternative hypothesis: highest value 934381231 is an outlier
```

No outlier detection test supports it.

Plot Box Office Revenues

```
df_summary <- df %>%
                group_by(Gender) %>%
                summarize(count = n(),
                          perc = n()/nrow(df),
                          mean = mean(AdjBoxOffice),
                          median = median(AdjBoxOffice))

color_m <- "#2980b9"
color_f <- "#27ae60"

print(df_summary)
```

```
Source: local data frame [2 x 5]

  Gender count      perc     mean   median
   (chr) (int)     (dbl)    (dbl)    (int)
1      f   467 0.2311881 65586882 44144648
2      m  1553 0.7688119 79786060 49841069
```

Plot Male and Female distributions separately, since medians are too close.

```
df_summary_m <- df_summary %>% filter(Gender=="m")

plot <- ggplot(df %>% filter(Gender=="m"), aes(x=AdjBoxOffice)) +
            geom_histogram(fill=color_m, bins=50, alpha=0.75) +
            fte_theme() +
            scale_x_log10(limits=c(10^7, 10^9), breaks=10^c(7:9), labels=c("$10M", "$100M",
                "$1B")) +
            geom_vline(xintercept=df_summary_m$mean, color="#1a1a1a") +
            geom_vline(xintercept=df_summary_m$median, color="#7f8c8d") +
            annotate(geom="text", label = "Mean:\n$79.8M", x=df_summary_f$mean+7*10^7, y=80,
                color="#1a1a1a", family="Source Sans Pro Bold", hjust=1, size=2) +
            annotate(geom="text", label = "Median:\n$49.8M", x=df_summary_m$median-0.5*10^7,
                y=80, color="#7f8c8d", family="Source Sans Pro Bold", hjust=1, size=2) +
            labs(title="Distribution of Box Office Revenues of Blockbusters w/ Male Lead",
                x="Domestic Box Office Revenue (2016 Dollars)", y="# of Movies")

max_save(plot, "movie-gender-1", "IMDb + Rotten Tomatoes via OMDb")

df_summary_f <- df_summary %>% filter(Gender=="f")

plot <- ggplot(df %>% filter(Gender=="f"), aes(x=AdjBoxOffice)) +
            geom_histogram(fill=color_f, bins=50, alpha=0.75) +
            fte_theme() +
            scale_x_log10(limits=c(10^7, 10^9), breaks=10^c(7:9), labels=c("$10M", "$100M",
                "$1B")) +
            geom_vline(xintercept=df_summary_f$mean, color="#1a1a1a") +
            geom_vline(xintercept=df_summary_f$median, color="#7f8c8d") +
            annotate(geom="text", label = "Mean:\n$65.6M", x=df_summary_f$mean+5*10^7, y=35,
                color="#1a1a1a", family="Source Sans Pro Bold", hjust=1, size=2) +
            annotate(geom="text", label = "Median:\n$44.1M", x=df_summary_f$median-0.5*10^7,
                y=35, color="#7f8c8d", family="Source Sans Pro Bold", hjust=1, size=2) +
            labs(title="Distribution of Box Office Revenues of Blockbusters w/ Female Lead",
                x="Domestic Box Office Revenue (2016 Dollars)", y="# of Movies")

max_save(plot, "movie-gender-2", "IMDb + Rotten Tomatoes via OMDb")
```

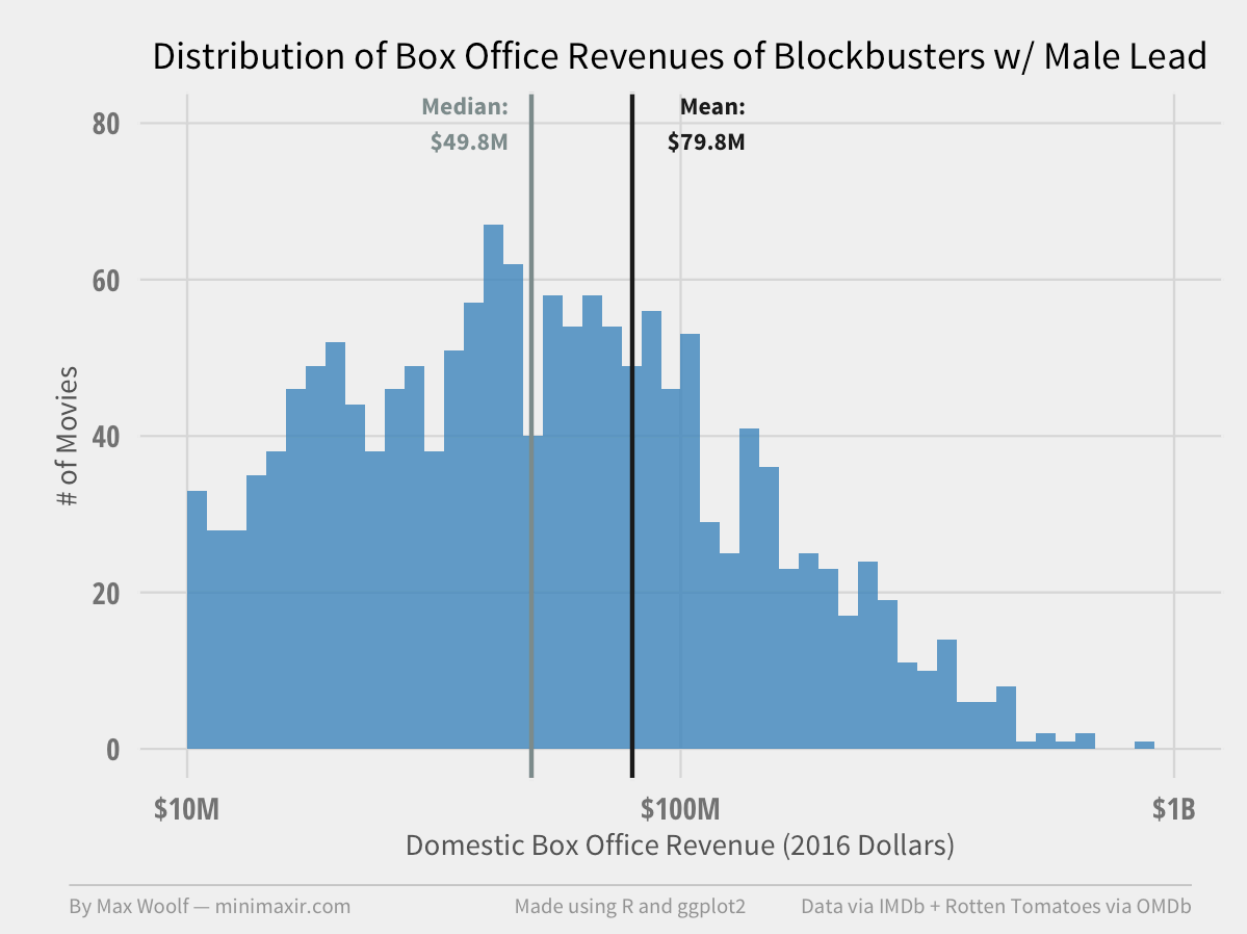
Distribution of Box Office Revenues of Blockbusters w/ Male Lead
Median:
$49.8M
Mean:
$79.8M
80
60
40
20
0
# of Movies
$10M
$100M
$1B
Domestic Box Office Revenue (2016 Dollars)
By Max Woolf — minimaxir.com
Made using R and ggplot2
Data via IMDb + Rotten Tomatoes via OMDb

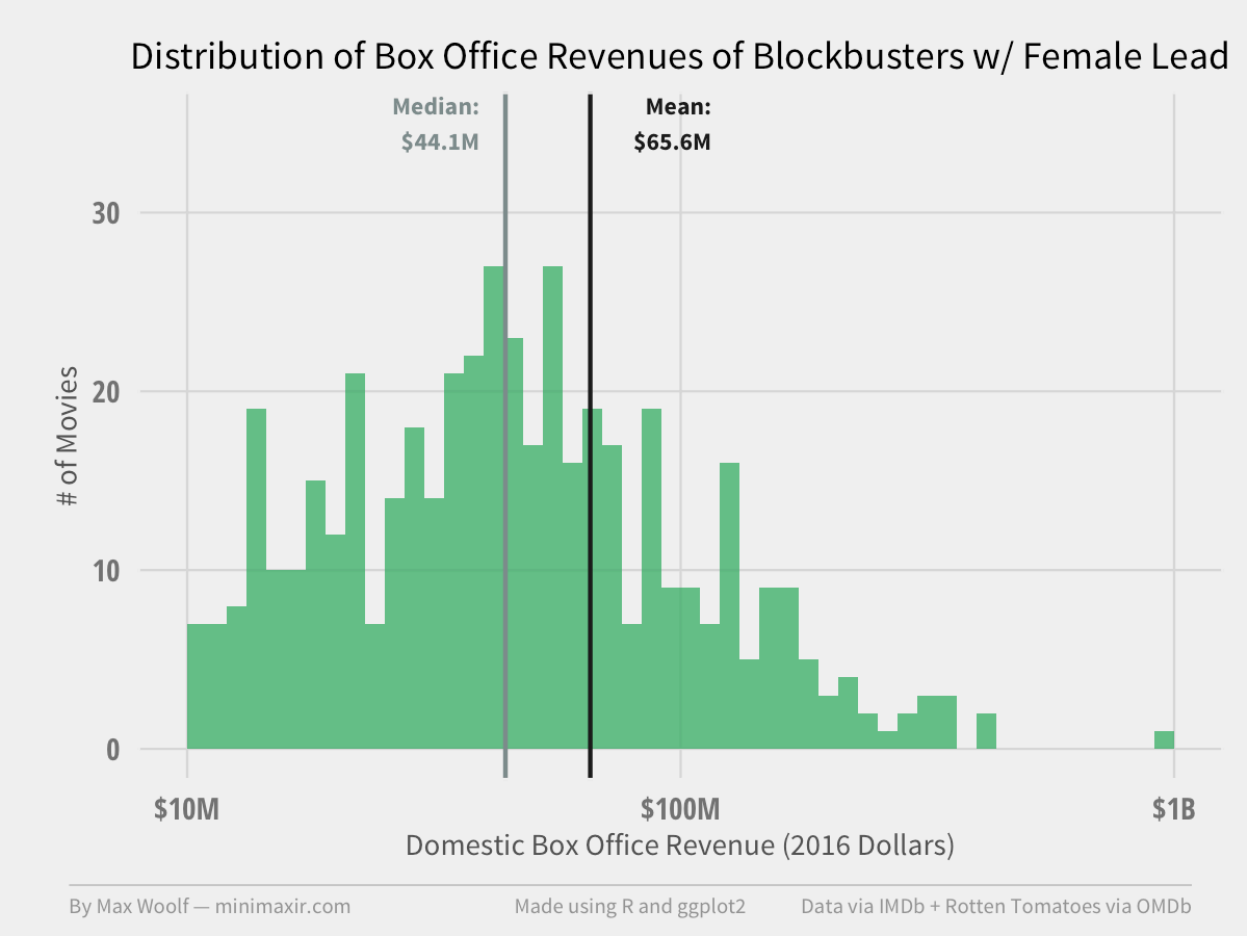

Plot kernel density distributions on each other.

```
plot <- ggplot(df, aes(x=AdjBoxOffice, fill=Gender)) +
            geom_density(alpha=0.75) +
            fte_theme() +
            scale_x_log10(limits=c(10^7, 10^9), breaks=10^c(7:9), labels=c("$10M", "$100M",
                "$1B")) +
            theme(legend.title = element_blank(), legend.position="top",
                legend.direction="horizontal", legend.key.width=unit(0.5, "cm"),
                legend.key.height=unit(0.25, "cm"), legend.margin=unit(0,"cm"),
                axis.title.y=element_blank(), axis.text.y=element_blank()) +
            scale_fill_manual(labels=c("Female Lead", "Male Lead"),
                values=c(color_f,color_m)) +
            labs(title="Density Distribution of B.O. Revenues of Blockbusters by Lead
                Gender", x="Domestic Box Office Revenue (2016 Dollars)")

max_save(plot, "movie-gender-3", "IMDb + Rotten Tomatoes via OMDb")
```

Do they come from a different distribution statistically speaking?

```
ks_test <- ks.test(
            unlist(df %>% filter(Gender=="m") %>% select(AdjBoxOffice)),
            unlist(df %>% filter(Gender=="f") %>% select(AdjBoxOffice)))

print(ks_test)
```

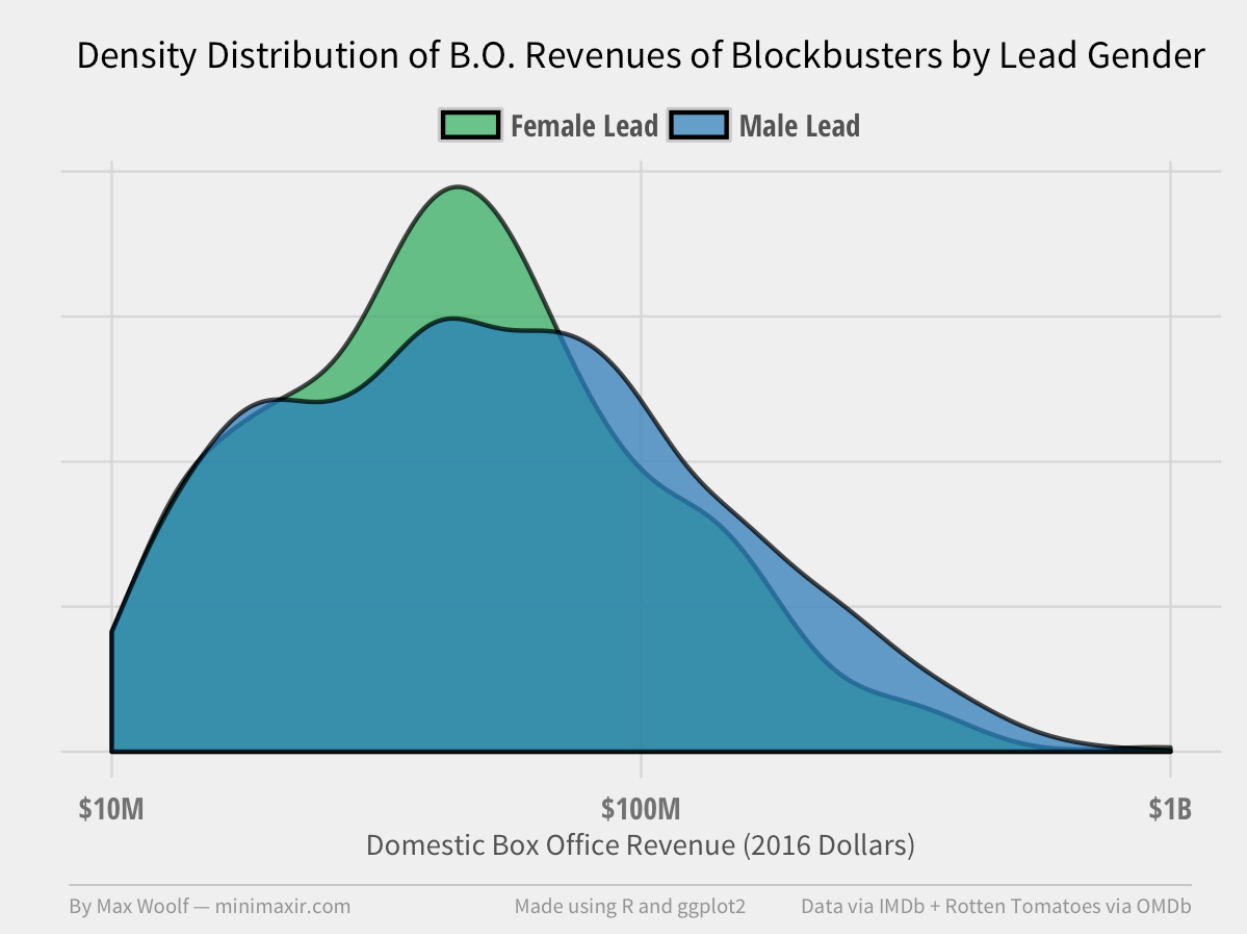
Density Distribution of B.O. Revenues of Blockbusters by Lead Gender
Female Lead
Male Lead
$10M
$100M
$1B
Domestic Box Office Revenue (2016 Dollars)
By Max Woolf — minimaxir.com
Made using R and ggplot2
Data via IMDb + Rotten Tomatoes via OMDb

Figure 1:

```r

## check if log-scaling changes the result

ks_test <- ks.test(
            log10(unlist(df %>% filter(Gender=="m") %>% select(AdjBoxOffice))),
            log10(unlist(df %>% filter(Gender=="f") %>% select(AdjBoxOffice))))

print(ks_test)
```

```
    Two-sample Kolmogorov-Smirnov test

data:  unlist(df %>% filter(Gender == "m") %>% select(AdjBoxOffice)) and unlist(df %>%
    filter(Gender == "f") %>% select(AdjBoxOffice))
D = 0.10585, p-value = 0.0006411
alternative hypothesis: two-sided


    Two-sample Kolmogorov-Smirnov test

data:  log10(unlist(df %>% filter(Gender == "m") %>% select(AdjBoxOffice))) and
    log10(unlist(df %>% filter(Gender == "f") %>% select(AdjBoxOffice)))
D = 0.10585, p-value = 0.0006411
alternative hypothesis: two-sided
```

The distributtion is different! Are the differences in means statistically significant?

```r
wilcox_test <- wilcox.test(
            unlist(df %>% filter(Gender=="m") %>% select(AdjBoxOffice)),
            unlist(df %>% filter(Gender=="f") %>% select(AdjBoxOffice)),
            alternative = "g")

print(wilcox_test)

## check if log-scaling changes the result

wilcox_test <- wilcox.test(
            log10(unlist(df %>% filter(Gender=="m") %>% select(AdjBoxOffice))),
            log10(unlist(df %>% filter(Gender=="f") %>% select(AdjBoxOffice))),
            alternative = "g")

print(wilcox_test)
```

```
    Wilcoxon rank sum test with continuity correction

data:  unlist(df %>% filter(Gender == "m") %>% select(AdjBoxOffice)) and unlist(df %>%
    filter(Gender == "f") %>% select(AdjBoxOffice))
W = 390070, p-value = 0.006514
alternative hypothesis: true location shift is greater than 0


    Wilcoxon rank sum test with continuity correction

data:  log10(unlist(df %>% filter(Gender == "m") %>% select(AdjBoxOffice))) and
    log10(unlist(df %>% filter(Gender == "f") %>% select(AdjBoxOffice)))
W = 390070, p-value = 0.006514
```

```
alternative hypothesis: true location shift is greater than 0
```

Plot Rotten Tomatoes Meter

Can reuse most of the code, unfortunately have to violate DRY for ad-hoc fixes.

```
df_summary <- df %>%
                group_by(Gender) %>%
                summarize(mean = mean(Meter, na.rm=T), median = median(Meter, na.rm=T))


print(df_summary)
```

```
Source: local data frame [2 x 3]

  Gender     mean median
   (chr)    (dbl)  (int)
1      f 47.97859     46
2      m 49.59381     49
```

```
df_summary_m <- df_summary %>% filter(Gender=="m")

plot <- ggplot(df %>% filter(Gender=="m"), aes(x=Meter)) +
            geom_histogram(fill=color_m, bins=50, alpha=0.75) +
            fte_theme() +
            scale_x_continuous(breaks=seq(0,100, by=10), limits=c(0, 100),
                labels=paste0(seq(0,100, by=10),"%")) +
            geom_vline(xintercept=df_summary_m$mean, color="#1a1a1a") +
            geom_vline(xintercept=df_summary_m$median, color="#7f8c8d") +
            annotate(geom="text", label = "Mean:\n49.6%", x=df_summary_m$mean+6, y=60,
                color="#1a1a1a", family="Source Sans Pro Bold", size=2) +
            annotate(geom="text", label = "Median:\n49%", x=df_summary_m$median-6, y=60,
                color="#7f8c8d", family="Source Sans Pro Bold", size=2) +
            labs(title="Distribution of RT Scores of Blockbusters w/ Male Lead", x="Rotten
                Tomatoes Tomatometer Score", y="# of Movies")

max_save(plot, "movie-gender-4", "IMDb + Rotten Tomatoes via OMDb")

df_summary_f <- df_summary %>% filter(Gender=="f")

plot <- ggplot(df %>% filter(Gender=="f"), aes(x=Meter)) +
            geom_histogram(fill=color_f, bins=50, alpha=0.75) +
            fte_theme() +
            scale_x_continuous(breaks=seq(0,100, by=10), limits=c(0, 100),
                labels=paste0(seq(0,100, by=10),"%")) +
            geom_vline(xintercept=df_summary_f$mean, color="#1a1a1a") +
            geom_vline(xintercept=df_summary_f$median, color="#7f8c8d") +
            annotate(geom="text", label = "Mean:\n48.0%", x=df_summary_f$mean+6, y=30,
                color="#1a1a1a", family="Source Sans Pro Bold", size=2) +
            annotate(geom="text", label = "Median:\n46%", x=df_summary_f$median-6, y=30,
                color="#7f8c8d", family="Source Sans Pro Bold",  size=2) +
            labs(title="Distribution of RT Scores of Blockbusters w/ Female Lead", x="Rotten
                Tomatoes Tomatometer Score", y="# of Movies")

max_save(plot, "movie-gender-5", "IMDb + Rotten Tomatoes via OMDb")
```

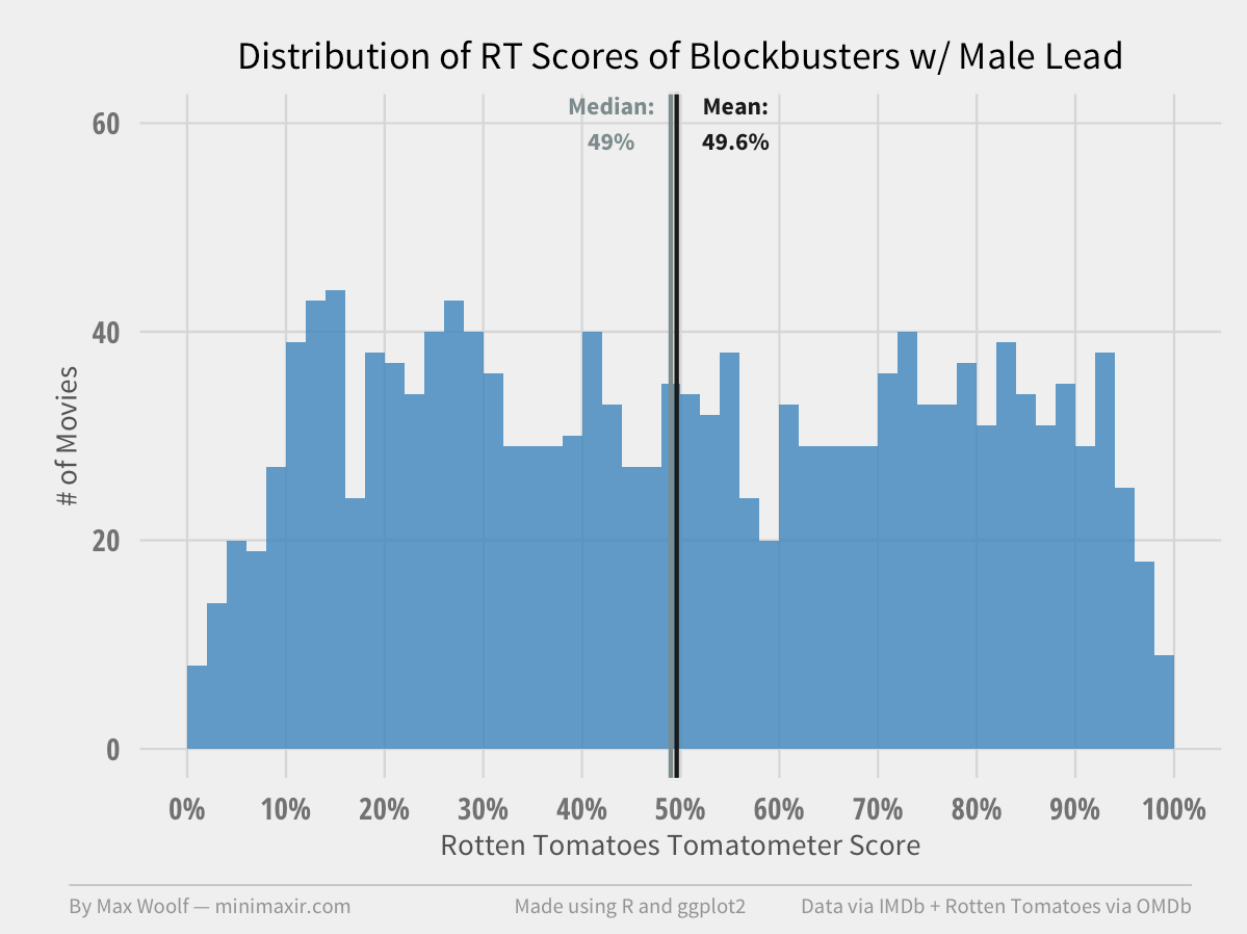
Distribution of RT Scores of Blockbusters w/ Male Lead
Median:
49%
Mean:
49.6%
60
40
20
0
# of Movies
0%
10%
20%
30%
40%
50%
60%
70%
80%
90%
100%
Rotten Tomatoes Tomatometer Score
By Max Woolf — minimaxir.com
Made using R and ggplot2
Data via IMDb + Rotten Tomatoes via OMDb

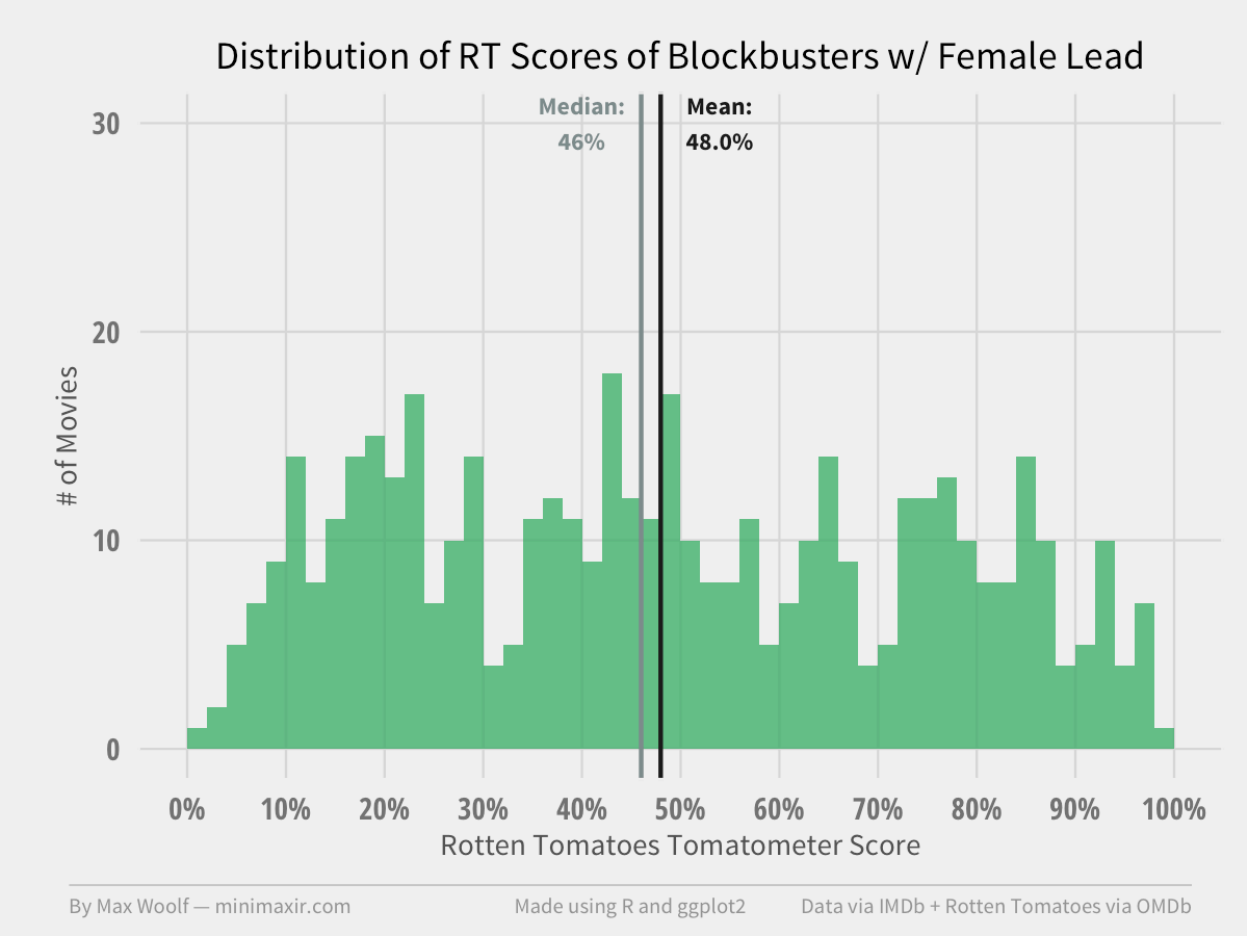


```
plot <- ggplot(df, aes(x=Meter, fill=Gender)) +
            geom_density(alpha=0.75) +
            fte_theme() +
            scale_x_continuous(breaks=seq(0,100, by=10), limits=c(0, 100),
                labels=paste0(seq(0,100, by=10),"%")) +
            theme(legend.title = element_blank(), legend.position="top",
                legend.direction="horizontal", legend.key.width=unit(0.5, "cm"),
                legend.key.height=unit(0.25, "cm"), legend.margin=unit(0,"cm"),
                axis.title.y=element_blank(), axis.text.y=element_blank()) +
            scale_fill_manual(labels=c("Female Lead", "Male Lead"),
                values=c(color_f,color_m)) +
            labs(title="Density Distribution of RT Scores of Blockbusters by Lead Gender",
                x="Rotten Tomatoes Tomatometer Score")

max_save(plot, "movie-gender-6", "IMDb + Rotten Tomatoes via OMDb")
```

```
ks_test <- ks.test(
            unlist(df %>% filter(Gender=="m") %>% select(Meter)),
            unlist(df %>% filter(Gender=="f") %>% select(Meter)))

print(ks_test)

wilcox_test <- wilcox.test(
            unlist(df %>% filter(Gender=="m") %>% select(Meter)),
```

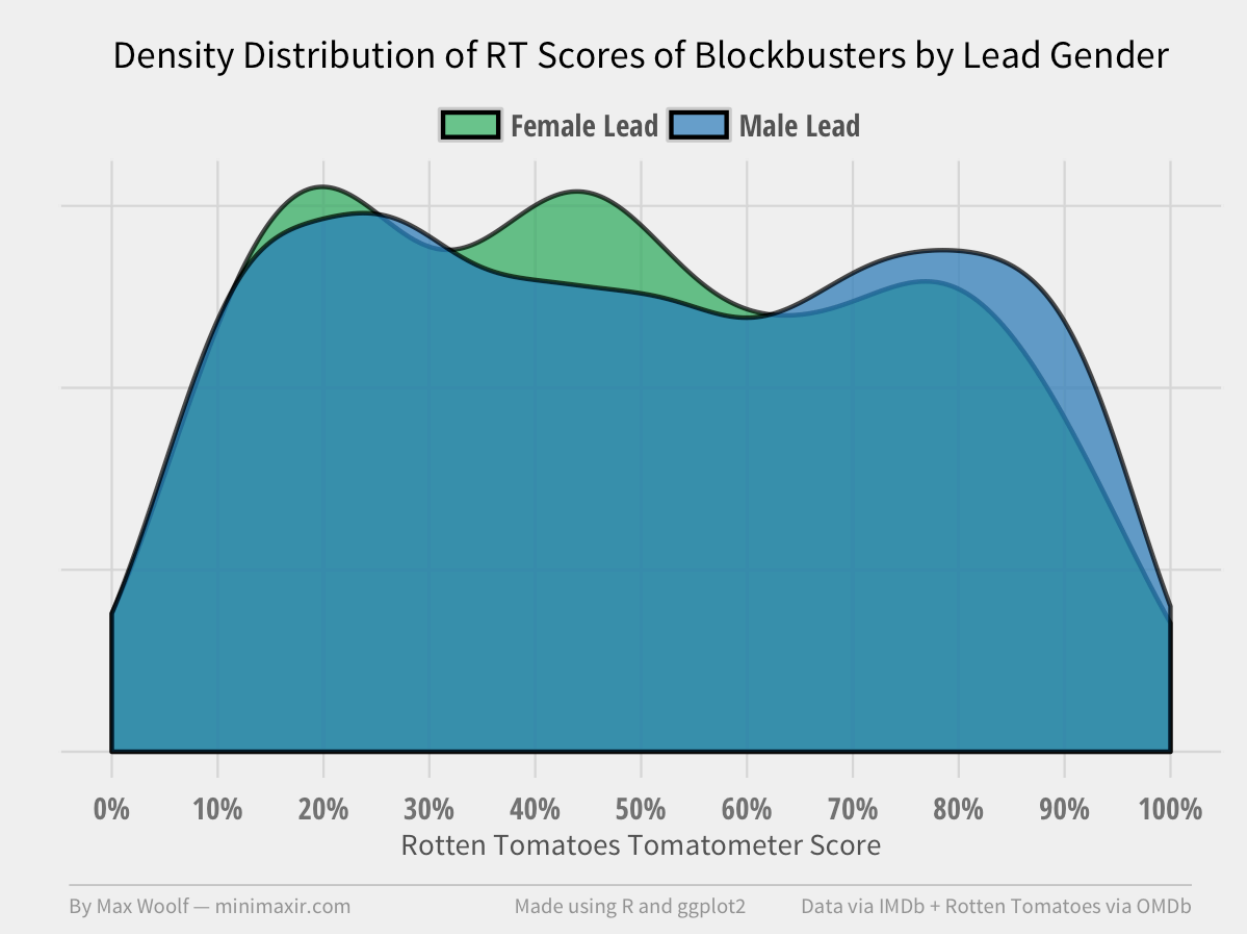
Density Distribution of RT Scores of Blockbusters by Lead Gender
Female Lead
Male Lead
0%
10%
20%
30%
40%
50%
60%
70%
80%
90%
100%
Rotten Tomatoes Tomatometer Score
By Max Woolf — minimaxir.com
Made using R and ggplot2
Data via IMDb + Rotten Tomatoes via OMDb

Figure 2:

```
            unlist(df %>% filter(Gender=="f") %>% select(Meter)),
            alternative="g")

print(wilcox_test)
```

```
    Two-sample Kolmogorov-Smirnov test

data:  unlist(df %>% filter(Gender == "m") %>% select(Meter)) and unlist(df %>%
    filter(Gender == "f") %>% select(Meter))
D = 0.048455, p-value = 0.3684
alternative hypothesis: two-sided


    Wilcoxon rank sum test with continuity correction

data:  unlist(df %>% filter(Gender == "m") %>% select(Meter)) and unlist(df %>%
    filter(Gender == "f") %>% select(Meter))
W = 374460, p-value = 0.1326
alternative hypothesis: true location shift is greater than 0
```

Plot Metacritic

```
df_summary <- df %>%
                group_by(Gender) %>%
                summarize(mean = mean(Metacritic, na.rm=T), median = median(Metacritic,
                    na.rm=T))


print(df_summary)
```

```
Source: local data frame [2 x 3]

  Gender     mean median
   (chr)    (dbl)  (dbl)
1      f 50.78523     50
2      m 51.76032     51
```

```
df_summary_m <- df_summary %>% filter(Gender=="m")

plot <- ggplot(df %>% filter(Gender=="m"), aes(x=Metacritic)) +
            geom_histogram(fill=color_m, bins=50, alpha=0.75) +
            fte_theme() +
            scale_x_continuous(breaks=seq(0,100, by=10), limits=c(0, 100)) +
            geom_vline(xintercept=df_summary_m$mean, color="#1a1a1a") +
            geom_vline(xintercept=df_summary_m$median, color="#7f8c8d") +
            annotate(geom="text", label = "Mean:\n51.8", x=df_summary_m$mean+6, y=80,
                color="#1a1a1a", family="Source Sans Pro Bold", size=2) +
            annotate(geom="text", label = "Median:\n51", x=df_summary_m$median-6, y=80,
                color="#7f8c8d", family="Source Sans Pro Bold", size=2) +
            labs(title="Distribution of Metacritic Scores of Blockbusters w/ Male Lead",
                x="Metacritic Score", y="# of Movies")

max_save(plot, "movie-gender-7", "IMDb + Rotten Tomatoes via OMDb")
```

```
14
15 df_summary_f <- df_summary %>% filter(Gender=="f")
16
17 plot <- ggplot(df %>% filter(Gender=="f"), aes(x=Metacritic)) +
18             geom_histogram(fill=color_f, bins=50, alpha=0.75) +
19             fte_theme() +
20             scale_x_continuous(breaks=seq(0,100, by=10), limits=c(0, 100)) +
21             geom_vline(xintercept=df_summary_f$mean, color="#1a1a1a") +
22             geom_vline(xintercept=df_summary_f$median, color="#7f8c8d") +
23             annotate(geom="text", label = "Mean:\n50.8", x=df_summary_f$mean+6, y=30,
                   color="#1a1a1a", family="Source Sans Pro Bold", size=2) +
24             annotate(geom="text", label = "Median:\n50", x=df_summary_f$median-6, y=30,
                   color="#7f8c8d", family="Source Sans Pro Bold",  size=2) +
25             labs(title="Distribution of Metacritic Scores of Blockbusters w/ Female Lead",
                   x="Metacritic Score", y="# of Movies")
26
27 max_save(plot, "movie-gender-8", "IMDb + Rotten Tomatoes via OMDb")
```

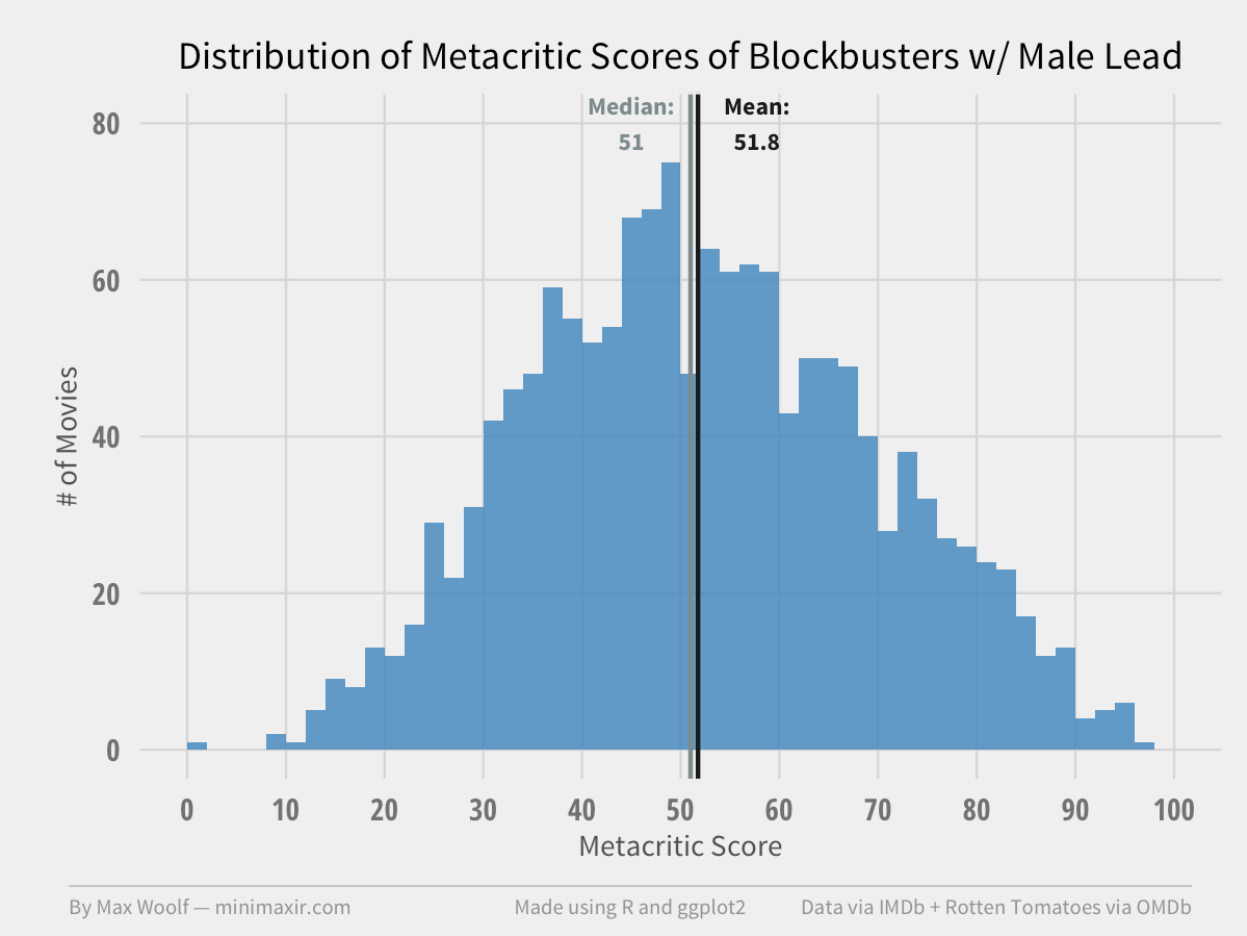

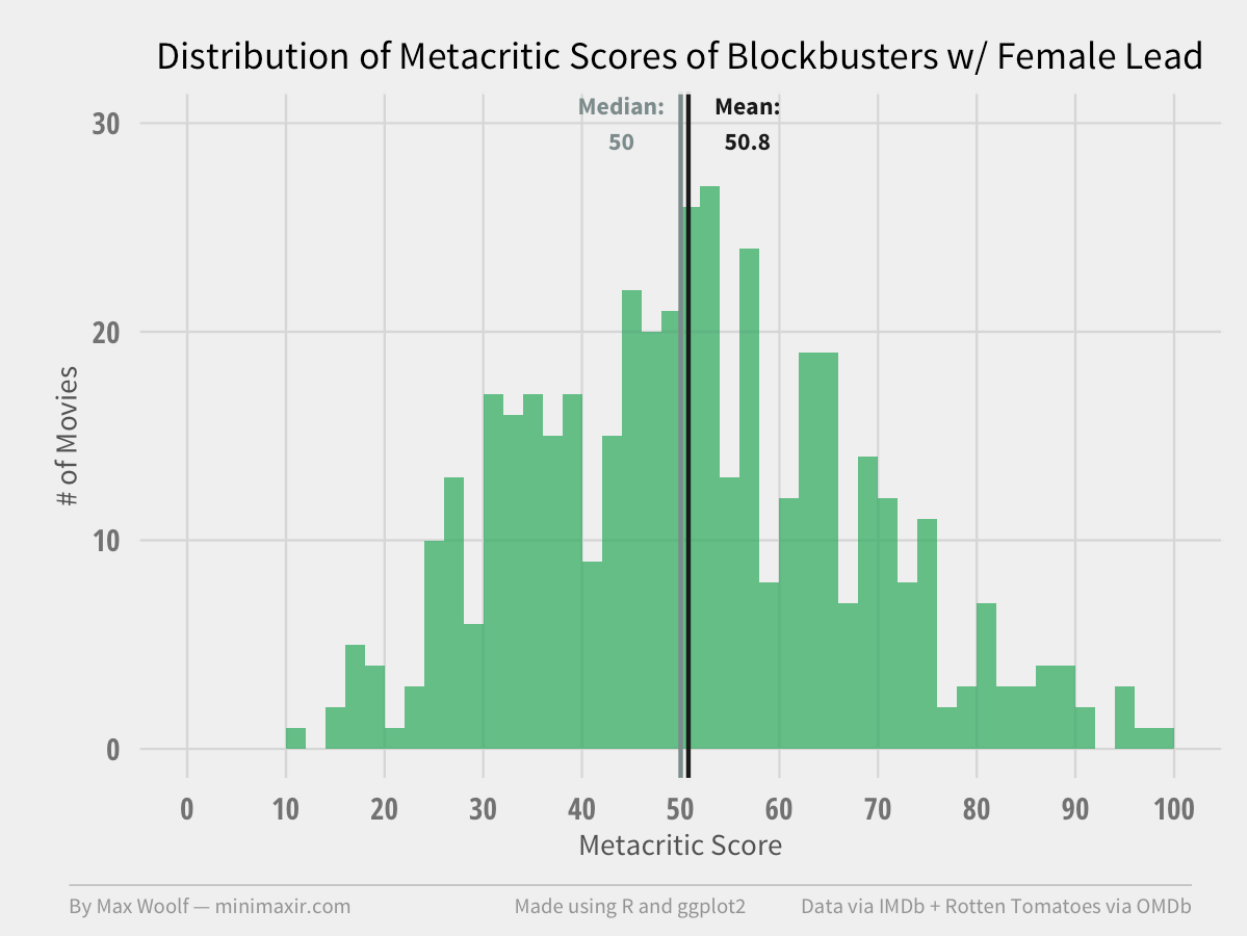


```
plot <- ggplot(df, aes(x=Metacritic, fill=Gender)) +
            geom_density(alpha=0.75) +
            fte_theme() +
            scale_x_continuous(breaks=seq(0,100, by=10), limits=c(0, 100)) +
            theme(legend.title = element_blank(), legend.position="top",
                legend.direction="horizontal", legend.key.width=unit(0.5, "cm"),
                legend.key.height=unit(0.25, "cm"), legend.margin=unit(0,"cm"),
                axis.title.y=element_blank(), axis.text.y=element_blank()) +
            scale_fill_manual(labels=c("Female Lead", "Male Lead"),
                values=c(color_f,color_m)) +
            labs(title="Density Distribution of Metacritic Scores of Blockbusters by Lead
                Gender", x="Metacritic Score")

max_save(plot, "movie-gender-9", "IMDb + Rotten Tomatoes via OMDb")
```

```
ks_test <- ks.test(
            unlist(df %>% filter(Gender=="m") %>% select(Metacritic)),
            unlist(df %>% filter(Gender=="f") %>% select(Metacritic)))

print(ks_test)

wilcox_test <- wilcox.test(
            unlist(df %>% filter(Gender=="m") %>% select(Metacritic)),
            unlist(df %>% filter(Gender=="f") %>% select(Metacritic)),
```

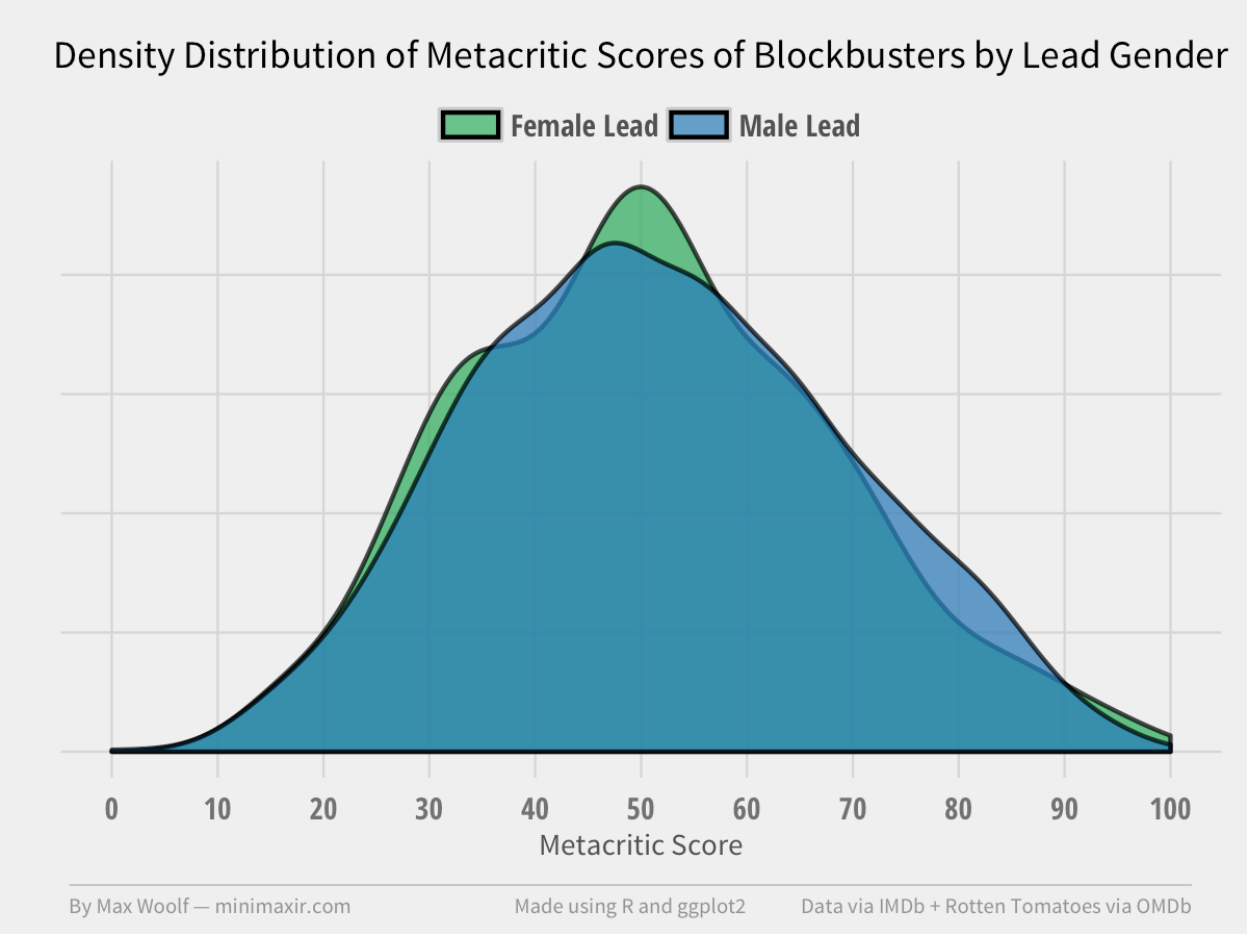
Density Distribution of Metacritic Scores of Blockbusters by Lead Gender
Female Lead
Male Lead
0
10
20
30
40
50
60
70
80
90
100
Metacritic Score
By Max Woolf — minimaxir.com
Made using R and ggplot2
Data via IMDb + Rotten Tomatoes via OMDb

Figure 3:

```
            alternative="g")

print(wilcox_test)
```

```
    Two-sample Kolmogorov-Smirnov test

data:  unlist(df %>% filter(Gender == "m") %>% select(Metacritic)) and unlist(df %>%
    filter(Gender == "f") %>% select(Metacritic))
D = 0.046268, p-value = 0.4521
alternative hypothesis: two-sided


    Wilcoxon rank sum test with continuity correction

data:  unlist(df %>% filter(Gender == "m") %>% select(Metacritic)) and unlist(df %>%
    filter(Gender == "f") %>% select(Metacritic))
W = 347130, p-value = 0.1368
alternative hypothesis: true location shift is greater than 0
```

Bootstramp Resample Means

```
resampleMeans <- function(df) {
    df_new <- df %>% sample_frac(replace=T)

    summary <- df_new %>%
                group_by(Gender) %>%
                summarize(AdjBoxOffice_m = mean(AdjBoxOffice),
                            Meter_m = mean(Meter, na.rm=T),
                            Metacritic_m = mean(Metacritic, na.rm=T))

    return (summary)
}

set.seed(4)
print(resampleMeans(df))
```

```
Source: local data frame [2 x 4]

  Gender AdjBoxOffice_m  Meter_m Metacritic_m
   (chr)          (dbl)    (dbl)        (dbl)
1      f       67986152 45.93206     49.40440
2      m       82138781 49.89780     52.02617
```

Pre-allocate space per this Stack Overflow answer.

```
resampleMovieData <- function(n) {

df_resample_summary <- data.frame(Gender = character(n*2), AdjBoxOffice_m = numeric(n*2),
                 Meter_m = numeric(n*2), Metacritic_m = numeric(n*2), stringsAsFactors =
                     FALSE)

for (i in seq(1,n*2 - 1, by = 2)) {
    df_resample_summary[c(i,i+1),] <- resampleMeans(df)
}

```

```
return(tbl_df(df_resample_summary))

}

set.seed(4)
print(resampleMovieData(4))
```

```
Source: local data frame [8 x 4]

  Gender AdjBoxOffice_m  Meter_m Metacritic_m
   (chr)          (dbl)    (dbl)        (dbl)
1      f       67986152 45.93206     49.40440
2      m       82138781 49.89780     52.02617
3      f       65405109 47.98039     51.02036
4      m       80331357 49.79551     51.76354
5      f       65073479 47.72557     50.11063
6      m       78003310 48.89669     51.36024
7      f       65424463 48.83369     51.77605
8      m       80139428 49.59100     51.92642
```

```
system.time( df_boot <- resampleMovieData(10000))

print(head(df_boot))
print(nrow(df_boot)) # expect 10000 * 2
```

```
   user  system elapsed
 37.460   4.230  42.968



Source: local data frame [6 x 4]

  Gender AdjBoxOffice_m  Meter_m Metacritic_m
   (chr)          (dbl)    (dbl)        (dbl)
1      f       67708627 48.58747     51.07289
2      m       79596707 50.01286     52.30313
3      f       70793578 47.57675     50.28372
4      m       77681742 50.53171     52.40521
5      f       72614163 47.77322     50.12472
6      m       78020424 49.21093     51.50498
[1] 20000
```

```
df_boot_agg <- df_boot %>%
                group_by(Gender) %>%
                summarize(
                    AdjBoxOffice_res_m = mean(AdjBoxOffice_m),
                    AdjBoxOffice_low_ci = quantile(AdjBoxOffice_m, 0.025),
                    AdjBoxOffice_high_ci = quantile(AdjBoxOffice_m, 0.975)
                        )

print(df_boot_agg)
```

```
Source: local data frame [2 x 4]

```

```
  Gender AdjBoxOffice_res_m AdjBoxOffice_low_ci AdjBoxOffice_high_ci
   (chr)              (dbl)               (dbl)                (dbl)
1      f           65531567            59238519             72485603
2      m           79758350            75590754             84168300
```

Plot Final Bootstrap

```
df_summary_means <- df %>% group_by(Gender) %>% summarize(mean = mean(AdjBoxOffice))

plot <- ggplot(df_boot, aes(x=AdjBoxOffice_m, fill=Gender)) +
    scale_x_continuous(limits=c(5*10^7, 10^8), breaks=seq(5*10^7,9*10^7, by=10^7),
        labels=paste0("$", seq(50,90, by=10), "M")) +
    scale_y_continuous(breaks=pretty_breaks(4)) +
    geom_histogram(bins=100, alpha=0.75, position="identity") +
    geom_point(mapping=aes(x=mean, y=0), data=df_summary_means, show.legend=F,
        color="black") +
    geom_errorbarh(mapping=aes(x=AdjBoxOffice_res_m, xmin=AdjBoxOffice_low_ci,
        xmax=AdjBoxOffice_high_ci, y=0), data=df_boot_agg, show.legend=F, color="black",
        height=0) +
    fte_theme() +
    theme(legend.title = element_blank(), legend.position="top",
        legend.direction="horizontal", legend.key.width=unit(0.5, "cm"),
        legend.key.height=unit(0.25, "cm"), legend.margin=unit(0,"cm")) +
    scale_fill_manual(labels=c("Female Lead", "Male Lead"), values=c(color_f,color_m)) +
    labs(title=sprintf("Resampled Avg. B.O. Revenues by Movie Lead Gender (n = %2d)",
        nrow(df_boot)/2), x="Average Domestic Box Office Revenue for Blockbusters (2016
        Dollars)", y="# of Resampled Averages")

max_save(plot, "movie-gender-10", "IMDb + Rotten Tomatoes via OMDb")
```

Determine P-Value of Final Bootstrap

Calculate the difference between the bootstrapped means; the P-value is the proportion of values where `m - f` < 0.

```
n <- 10000

means_vector <- unlist(df_boot$AdjBoxOffice)
means_diff <- c()

for (i in seq(1,n*2 - 1, by = 2)) {
    means_diff <- c(means_diff, means_vector[i+1] - means_vector[i])
}

print(means_diff[1:4])

print(sum(means_diff <= 0)/n)   # p-value of difference
```

```
[1] 11888080  6888164  5406261  8136011
[1] 2e-04
```

Bootstrap Movie!

Render each frame of the resample; composite into GIF later.

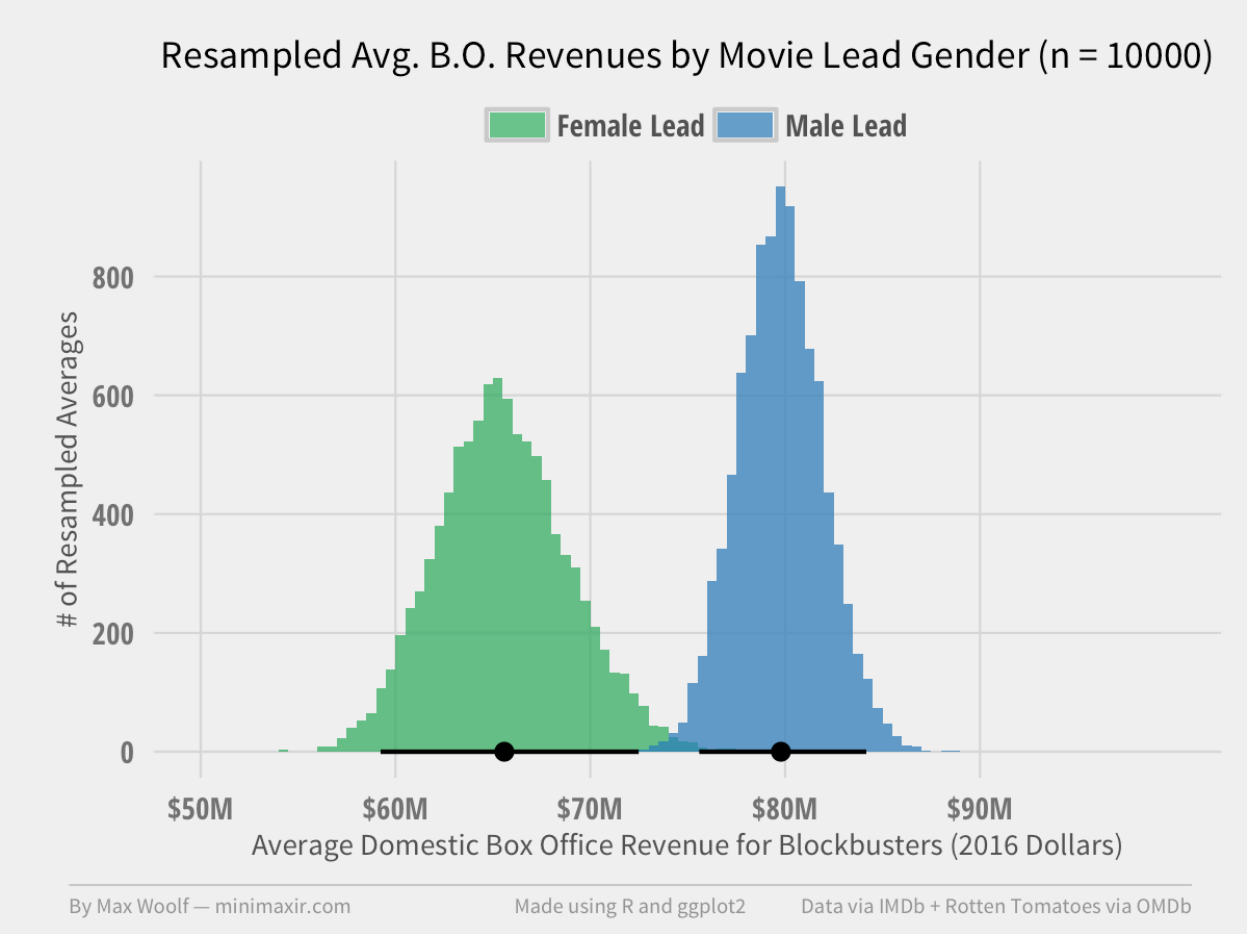
Resampled Avg. B.O. Revenues by Movie Lead Gender (n = 10000)
Female Lead
Male Lead
# of Resampled Averages
800
600
400
200
0
$50M
$60M
$70M
$80M
$90M
Average Domestic Box Office Revenue for Blockbusters (2016 Dollars)
By Max Woolf — minimaxir.com
Made using R and ggplot2
Data via IMDb + Rotten Tomatoes via OMDb

Figure 4:

```
system("mkdir -p movie_frames")

movie_frames <- function(size) {
    df_boot_sub <- df_boot %>% head(size*2)

    df_boot_agg_sub <- df_boot_sub %>%
                group_by(Gender) %>%
                summarize(
                    AdjBoxOffice_res_m = mean(AdjBoxOffice_m),
                    AdjBoxOffice_low_ci = quantile(AdjBoxOffice_m, 0.025),
                    AdjBoxOffice_high_ci = quantile(AdjBoxOffice_m, 0.975)
                        )

    plot <- ggplot(df_boot_sub, aes(x=AdjBoxOffice_m, fill=Gender)) +
    scale_x_continuous(limits=c(5*10^7, 10^8), breaks=seq(5*10^7,9*10^7, by=10^7),
        labels=paste0("$", seq(50,90, by=10), "M")) +
    scale_y_continuous(breaks=pretty_breaks(4)) +
    geom_histogram(bins=100, alpha=0.75, position="identity") +
    geom_point(mapping=aes(x=mean, y=0), data=df_summary_means, show.legend=F,
        color="black") +
    geom_errorbarh(mapping=aes(x=AdjBoxOffice_res_m, xmin=AdjBoxOffice_low_ci,
        xmax=AdjBoxOffice_high_ci, y=0), data=df_boot_agg_sub, show.legend=F, color="black",
        height=0) +
    fte_theme() +
    theme(legend.title = element_blank(), legend.position="top",
        legend.direction="horizontal", legend.key.width=unit(0.5, "cm"),
        legend.key.height=unit(0.25, "cm"), legend.margin=unit(0,"cm")) +
    scale_fill_manual(labels=c("Female Lead", "Male Lead"), values=c(color_f,color_m)) +
    labs(title=sprintf("Resampled Avg. B.O. Revenues by Movie Lead Gender (n = %2d)", size),
        x="Average Domestic Box Office Revenue for Blockbusters (2016 Dollars)", y="# of
        Resampled Averages")

max_save(plot, sprintf("movie_frames/movie_%06d", size), "IMDb + Rotten Tomatoes via OMDb")

}
```

```
system.time( x <- lapply(seq(100,10000,100), movie_frames) )
```

```
   user  system elapsed
 54.031   2.565  60.278
```

The MIT License (MIT)

NESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.